

Advanced Certification in

Applied Data Science, Machine Learning & IoT

By E&ICT Academy, IIT Guwahati

Project: Hotel Reviews Sentiment Analysis

Name: [Rajesh Bisht](#)

Email: rbisht.india@gmail.com

Batch: 8 (Eight)

Git Link: [RajeshBisht28/IIT_ML_HotelReviewSentiment](#)

Description: ML project for Hotel Review Sentiment Analysis.

In []:

In []:



Hotel Reviews sentiment analysis.



Dataset - Column Explanation:

The dataset contains hotel reviews with the following columns:

Column Name	Description
Hotel_Address	Address of the hotel.
Additional_Number_of_Scoring	Additional number of reviews for the hotel.
Review_Date	Date when the review was posted.
Average_Score	Average score of the hotel based on all reviews.
Hotel_Name	Name of the hotel.
Reviewer_Nationality	Nationality of the reviewer.
Negative_Review	Text of the negative review (if any).
Review_Total_Negative_Word_Counts	Word count of the negative review.
Total_Number_of_Reviews	Total number of reviews for the hotel.
Positive_Review	Text of the positive review (if any).
Review_Total_Positive_Word_Counts	Word count of the positive review.
Total_Number_of_Reviews_Reviewer_Has_Given	Total reviews the reviewer has given across hotels.
Reviewer_Score	Score given by the reviewer.
Tags	Tags related to the review (e.g., trip type, room type).

Column Name	Description
days_since_review	Days since the review was posted.
lat	Latitude of the hotel.
lng	Longitude of the hotel.

✗ Columns NOT Important for Sentiment Analysis:

These columns do not directly impact sentiment classification:

- 🏠 **Hotel_Address** – Location is not relevant.
- 📊 **Additional_Number_of_Scoring** – Unrelated to individual reviews.
- 📅 **Review_Date** – Date does not influence sentiment.
- ⭐ **Average_Score** – A general rating, not review-specific.
- 🏨 **Hotel_Name** – Hotel name does not impact sentiment.
- 🌐 **Reviewer_Nationality** – Nationality is not crucial for sentiment.
- 📋 **Total_Number_of_Reviews** – Does not affect individual sentiment.
- 🔄 **Total_Number_of_Reviews_Reviewer_Has_Given** – Unrelated to review sentiment.
- 🏷️ **Tags** – Useful for filtering, not sentiment.
- ⌚ **days_since_review** – Time does not affect sentiment directly.
- 📍 **lat, lng** – Location coordinates are irrelevant.

✓ Important Columns for Sentiment Analysis:

These columns are key for extracting sentiment:

- 📄 **Negative_Review** – Essential for detecting negative sentiment.
- 😊 **Positive_Review** – Essential for detecting positive sentiment.
- 📋 **Review_Total_Negative_Word_Counts** – Indicates negativity intensity.
- 📋 **Review_Total_Positive_Word_Counts** – Indicates positivity intensity.

- 🌟 **Reviewer_Score** – Can be used to validate sentiment classification.

Overview

Tool: Sentiment Analysis with Python

Sentiment Analysis?

Sentiment analysis is a crucial part of **Natural Language Processing (NLP)**. It involves **extracting emotions** from raw text to determine whether a statement conveys a **positive, negative, or neutral** sentiment.



This technique is widely used for:

- 🗣️ **Social media monitoring** – Understanding public opinion.
- 🌟 **Customer reviews analysis** – Measuring customer satisfaction.
- 📊 **Brand reputation tracking** – Identifying trends in feedback.

The goal of this study is to demonstrate how sentiment analysis can be performed using **Python**.

Libraries We Will Use:

For this task, we will leverage the following key libraries:

 Library	 Purpose
NLTK	The most popular Python library for NLP techniques. POS, tokenizer, lemmatizer, wordnet.
Gensim	A toolkit for topic modeling and vector space modeling.
Scikit-learn	The most widely used machine learning library in Python.



Dataset: Hotel Reviews

We will analyze a dataset containing **hotel reviews** from various customers. Each observation consists of:

- **A textual review** – The customer's experience in their own words.
- **An overall rating** – A numerical score given by the customer.



Dataset Source:



[515K Hotel Reviews Data \(Europe\)](#)



Problem Statement

For each customer review, our goal is to predict whether it is **positive (good experience)** or **negative (bad experience)** based only on the raw textual feedback.



Sentiment Classification:

We will categorize reviews into **two sentiment classes** based on their overall ratings:

- 👎 **Negative Reviews** → Ratings < 5
- 👍 **Positive Reviews** → Ratings ≥ 5



Loading the Raw Data

Let's sentiment analysis, we first **load the raw dataset** containing hotel reviews.



Data Structure:

Each **textual review** is divided into two parts:

- **Positive Review** ✨ – Highlights what the customer liked.

- **Negative Review** ❌ – Mentions aspects that the customer disliked.

🔧 Data Preprocessing:

To simplify our analysis, we will:

1. **Merge** both the positive and negative review sections.
2. **Create a single text column** containing the full customer feedback.
3. **Remove unnecessary metadata** to focus only on the raw text data.

By doing this, we ensure that our analysis is based purely on textual content without any additional numerical scores or labels.

```
In [1]: import pandas as pd
# Load data from CSV
main_df = pd.read_csv("Hotel_Reviews.csv")
# Concatenate neative and positive reviews.
main_df["review"] = main_df["Negative_Review"] + main_df["Positive_Review"]
# create the label
main_df["is_bad_review"] = main_df["Reviewer_Score"].apply(lambda x: 1 if x < 5 else 0)
# select only relevant columns
main_df = main_df[["review", "is_bad_review"]]
main_df.head()
```

```
Out[1]:
```

	review	is_bad_review
0	I am so angry that i made this post available...	1
1	No Negative No real complaints the hotel was g...	0
2	Rooms are nice but for elderly a bit difficul...	0
3	My room was dirty and I was afraid to walk ba...	1
4	You When I booked with your company on line y...	0

Data set Rows, Columns Counts.

```
In [2]: no_rows, no_cols = main_df.shape
        print(f"Rows Counts : {no_rows} , Columns Counts: {no_cols}")
```

Rows Counts : 515738 , Columns Counts: 2



Sample Data ~ 10% Of Dataset for reliable computation

```
In [4]: sample_df = main_df.sample(frac = 0.1, replace = False, random_state=42)
        no_rows, no_cols = sample_df.shape
        print(f"Rows Counts : {no_rows} , Columns Counts: {no_cols}")
```

Rows Counts : 51574 , Columns Counts: 2

```
In [5]: sample_df.isnull().sum()
```

```
Out[5]: review          0
        is_bad_review    0
        dtype: int64
```

No Null values exist in sample data



Text Preprocessing: Cleaning the Data



Removing Default Placeholder Comments

In our dataset, customers who did not leave specific feedback have the following placeholders:

- **"No Negative"** → Appears when no negative feedback is provided.
- **"No Positive"** → Appears when no positive feedback is given.

These placeholders do not carry any real sentiment and **must be removed** from our text data.



Cleaning Text Data

Once we remove placeholder comments, the next step involves **text preprocessing** to ensure better accuracy in sentiment analysis.

We will perform the following operations:

- ✓ **Lowercasing** – Convert all text to lowercase for uniformity.
- ✓ **Removing Punctuation** – Eliminate special characters and symbols.
- ✓ **Tokenization** – Split sentences into individual words (tokens).
- ✓ **Stopword Removal** – Remove common words like "the", "is", "and" that do not add sentiment value.
- ✓ **Lemmatization/Stemming** – Reduce words to their base form (e.g., "running" → "run").

These steps are crucial for improving our **Natural Language Processing (NLP) model's performance** and ensuring more accurate sentiment predictions.

🚀 **Let's dive into the implementation!**

Preprocess Text process :: for sample_df , all basic processing of NLP

```
In [6]: # return the wordnet values according to the POS tag
from nltk.corpus import wordnet

def fetch_wordnet_pos(pos_tag):
    if pos_tag.startswith('J'):
        return wordnet.ADJ
    elif pos_tag.startswith('V'):
        return wordnet.VERB
    elif pos_tag.startswith('N'):
        return wordnet.NOUN
    elif pos_tag.startswith('R'):
        return wordnet.ADV
    else:
        return wordnet.NOUN

import string
from nltk import pos_tag
from nltk.corpus import stopwords
from nltk.tokenize import WhitespaceTokenizer
from nltk.stem import WordNetLemmatizer
```



```
def preprocess_text(text):
    # lower text
    text = text.lower()
    # tokenize text and remove puncutation
    text = [word.strip(string.punctuation) for word in text.split(" ")]
    # remove words that contain numbers
    text = [word for word in text if not any(c.isdigit() for c in word)]
    # remove stop words
    stop = stopwords.words('english')
    text = [x for x in text if x not in stop]
    # remove empty tokens
    text = [t for t in text if len(t) > 0]
    # find pos tag text
    pos_tags = pos_tag(text)
    # Lemmatization on a list of part-of-speech (POS) tagged words
    text = [WordNetLemmatizer().lemmatize(t[0], fetch_wordnet_pos(t[1])) for t in pos_tags]
    # remove words with only one letter
    text = [t for t in text if len(t) > 1]
    # join together
    text = " ".join(text)
    return(text)
```

```
In [8]: # remove 'No Negative' or 'No Positive' from text
sample_df["review"] = sample_df["review"].apply(lambda x: x.replace("No Negative", "").replace("No Positive", ""))
```

```
In [9]: ### Preprocess sample data
sample_df["review_clean"] = sample_df["review"].apply(lambda x: preprocess_text(x))
```

✨ Feature engineering

Adding Sentiment Analysis to Reviews

📌 Overview

To analyze customer feedback, we use **Sentiment Analysis**, which helps classify reviews as **positive, negative, or neutral** based on their text content.

In this step, we utilize **NLTK's SentimentIntensityAnalyzer (VADER)** to compute sentiment scores for each review.

```
In [10]: # Import the VADER Sentiment Analyzer from NLTK
from nltk.sentiment.vader import SentimentIntensityAnalyzer

# Initialize Sentiment Intensity Analyzer
sid = SentimentIntensityAnalyzer()

# Apply Sentiment Analysis to each review text
sample_df["sentiments"] = sample_df["review"].apply(lambda x: sid.polarity_scores(x))

# Expand sentiment dictionary into separate columns and merge with the main DataFrame
sample_df = pd.concat([sample_df.drop(['sentiments'], axis=1), sample_df['sentiments'].apply(pd.Series)], axis=1)

In [11]: reviews_df = sample_df.copy()
```

Creating Doc2Vec Vector Columns for Text Data

```
In [14]: ## Creating Doc2Vec Vector Columns for Text Data
from gensim.test.utils import common_texts
## This code transforms text data into numerical vectors using Doc2Vec from the gensim Library.
from gensim.models.doc2vec import Doc2Vec, TaggedDocument
## Convert Text Data into TaggedDocument Format
documents = [TaggedDocument(doc, [i]) for i, doc in enumerate(reviews_df["review_clean"].apply(lambda x: x.split(" ")))
# train a Doc2Vec model
model = Doc2Vec(documents, vector_size=5, window=2, min_count=1, workers=4)
### Convert Each Document into a Vector
doc2vec_df = reviews_df["review_clean"].apply(lambda x: model.infer_vector(x.split(" "))).apply(pd.Series)
## Merge the Vector Data Back into the Original DataFrame
doc2vec_df.columns = ["doc2vec_vector_" + str(x) for x in doc2vec_df.columns]
```

Doc2Vec Generated Vectors Table

doc2vec_vector_0	doc2vec_vector_1	doc2vec_vector_2	doc2vec_vector_3	doc2vec_vector_4
0.12	-0.45	0.78	0.34	-0.23

-0.21 0.67 -0.88 0.13 0.45

Tagged Documents Representation

```
[ TaggedDocument(words=['great', 'product', 'love', 'it'], tags=[0]),  
  TaggedDocument(words=['poor', 'quality', 'waste', 'money'], tags=[1])  
]
```

**Example - What we are getting now:

This **merges** the original DataFrame (`reviews_df`) with the new DataFrame (`doc2vec_df`) containing **vectorized text representations**.

1 Original `reviews_df`

review_clean
"great product love it"
"poor quality waste money"

2 `doc2vec_df` (Generated Vectors)

doc2vec_vector_0	doc2vec_vector_1	doc2vec_vector_2	doc2vec_vector_3	doc2vec_vector_4
0.12	-0.45	0.78	0.34	-0.23
-0.21	0.67	-0.88	0.13	0.45

3 After `pd.concat()` (Final `reviews_df`)

review_clean	doc2vec_vector_0	doc2vec_vector_1	doc2vec_vector_2	doc2vec_vector_3	doc2vec_vector_4
"great product love it"	0.12	-0.45	0.78	0.34	-0.23

review_clean	doc2vec_vector_0	doc2vec_vector_1	doc2vec_vector_2	doc2vec_vector_3	doc2vec_vector_4
"poor quality waste money"	-0.21	0.67	-0.88	0.13	0.45

Now, each text review has **5 numerical vector values** representing its semantic meaning.

- ✓ Merges the **original text data** with the **generated document vectors**
- ✓ Allows **further processing**, such as training Machine Learning models
- ✓ Keeps the **original review data intact** while adding meaningful features

Extracts TF-IDF (Term Frequency-Inverse Document Frequency) features from the review_clean column.

In []:

```
In [15]: from sklearn.feature_extraction.text import TfidfVectorizer
import pandas as pd

# Sample DataFrame
###reviews_df = pd.DataFrame({"review_clean": ["This is a sample review", "Another review text", "More text here"]})

# Apply TF-IDF
tfidf = TfidfVectorizer(min_df=10)
tfidf_result = tfidf.fit_transform(reviews_df["review_clean"]).toarray()

# Create DataFrame with TF-IDF values
tfidf_df = pd.DataFrame(tfidf_result, columns=tfidf.get_feature_names_out())
tfidf_df.columns = ["word_" + str(x) for x in tfidf_df.columns]
tfidf_df.index = reviews_df.index

# Concatenate with original DataFrame
reviews_df = pd.concat([reviews_df, tfidf_df], axis=1)
```

The TF-IDF metric solves this problem:

- TF computes the classic number of times the word appears in the text
- IDF computes the relative importance of this word which depends on how many texts the word can be found

◆ Example: Before & After TF-IDF Transformation

 Original `reviews_df`

review_clean
"great product love it"
"poor quality waste money"

 After Adding TF-IDF Features

review_clean	word_great	word_product	word_love	word_quality	word_money
"great product love it"	0.45	0.38	0.52	0.00	0.00
"poor quality waste money"	0.00	0.00	0.00	0.58	0.47

Each review is now represented by **TF-IDF scores**, making it suitable for **Machine Learning tasks** such as sentiment analysis and classification.

In [32]: `reviews_df.head()`

Out[32]:

	review	is_bad_review	review_clean	neg	neu	pos	compound	doc2vec_vector_0	doc2vec_vector_1	doc2vec_
488440	Would have appreciated a shop in the hotel th...	0	would appreciate shop hotel sell drinking wate...	0.049	0.617	0.334	0.9924	0.057766	0.483825	-(
274649	No tissue paper box was present at the room	0	tissue paper box present room	0.216	0.784	0.000	-0.2960	0.072940	-0.038417	-(
374688	Pillows Nice welcoming and service	0	pillow nice welcome service	0.000	0.345	0.655	0.6908	-0.093065	0.059668	(
404352	Everything including the nice upgrade The Hot...	0	everything include nice upgrade hotel revamp s...	0.000	0.621	0.379	0.9153	-0.124238	0.099710	(
451596	Lovely hotel v welcoming staff	0	lovely hotel welcome staff	0.000	0.230	0.770	0.7717	0.089844	0.108858	(



In [16]: `reviews_df.shape`

Out[16]: (51574, 3833)

Exploratory data analysis

Let's try understanding our data, what it contains and how they concatenate.

```
In [17]: # show is_bad_review distribution
reviews_df["is_bad_review"].value_counts(normalize = True)
```

```
Out[17]: is_bad_review
0      0.956761
1      0.043239
Name: proportion, dtype: float64
```

Our dataset is highly imbalanced because less than 5% of our reviews are considered as negative ones. This information will be very useful for the modelling part.

Wordclouds ==> For Reviews:

```
In [18]: # wordcloud presentation
# total_chars ==> total characters in the characters column
reviews_df["total_chars"] = reviews_df["review"].apply(lambda x: len(x))

### total words in the words column
reviews_df["total_words"] = reviews_df["review"].apply(lambda x: len(x.split(" ")))

from wordcloud import WordCloud
import matplotlib.pyplot as plt

def show_wordcloud(data, title = None):
    wordcloud = WordCloud(
        background_color = 'white',
        max_words = 200,
        max_font_size = 40,
        scale = 3,
        random_state = 42
    ).generate(str(data))

    fig = plt.figure(1, figsize = (20,20))
    plt.axis('off')
    if title:
        fig.suptitle(title, fontsize = 15)
        fig.subplots_adjust(top = 2.3)

    plt.imshow(wordcloud)
```

[illegible]

```
# highest positive sentiment reviews (with more than 5 words)
reviews_df[reviews_df["total_words"] >= 5].sort_values("pos", ascending = False)[["review", "pos"]].head(10)
```


Out[19]:

	review	pos
43101	A perfect location comfortable great value	0.931
211742	Clean comfortable lovely staff	0.907
175551	Friendly welcome Comfortable room	0.905
365085	Good location great value	0.904
109564	Clean friendly and comfortable	0.902
145743	Good value amazing location	0.901
407590	breakfast excellent Clean comfort	0.899
407546	Great place I enjoyed	0.881
218571	Beautiful Quirky Comfortable	0.878
436901	Lovely comfortable rooms	0.877

- **Key Finding:** The most positive reviews correspond to positive customer feedback.
- **Implication:** This suggests a direct link between review sentiment and customer satisfaction.

```
In [20]: # Lowest negative sentiment reviews (with more than 5 words)
reviews_df[reviews_df["total_words"] >= 5].sort_values("neg", ascending = False)[["review", "neg"]].head(10)
```

Out[20]:

		review	neg
193086	No dislikes	LOCATION	0.831
356368	Nothing Great	helpful wonderful staff	0.812
318516	A disaster	Nothing	0.804
458794	Nothing Excellent	friendly helpful staff	0.799
29666	A bit noisy	No	0.796
426057	Dirty hotel	Smells bad	0.762
263187	Very bad service	No	0.758
443796	Nothing perfect		0.750
181508	Window blind was broken		0.744
175316	Nothing Super	friendly staff	0.743


- **VADER's Limitations:**
 - Phrases like "no problems" or "nothing wrong" can be misclassified.
- **Overall Trend:**
 - Most negative reviews accurately reflect negative customer sentiment.

In [21]: *# plot sentiment distribution for positive and negative reviews*

```
import seaborn as sns

for x in [0, 1]:
    subset = reviews_df[reviews_df['is_bad_review'] == x]

    # Draw the density plot
    if x == 0:
        label = "Good reviews"
    else:
        label = "Bad reviews"
    sns.kdeplot(subset['compound'], label=label)
```



Sentiment Distribution Analysis

The above graph illustrates the **distribution of sentiment scores** among **good and bad reviews**.

- **Good reviews** are mostly classified as **very positive** by the Vader sentiment analysis tool.
- **Bad reviews** tend to have **lower compound sentiment scores**, indicating more negative sentiment.

◆ Key Insight:

This analysis confirms that the **computed sentiment features** will play a crucial role in the **modeling phase**, helping to distinguish between positive and negative reviews effectively. 🚀

Feature Selection & Train-Test Split

Let's perform **feature selection and dataset splitting** for training a machine learning model.

◆ Steps Explained:

1. 🔍 Feature Selection:

- The **target label** (`is_bad_review`), **raw text** (`review`), and **cleaned text** (`review_clean`) are **excluded** from the feature set.
- Remaining columns in `reviews_df` are selected as **features** for model training.

2. 📊 Train-Test Split:

- The dataset is **split into training (80%) and testing (20%) subsets**.
- `train_test_split()` ensures **randomized sampling** with a fixed seed (`random_state=42`) for reproducibility.


🚀 Purpose:

This step **prepares the dataset** for training a **Random Forest Classifier** by selecting meaningful features and ensuring a proper training/testing distribution.

```
In [22]: # feature selection
label = "is_bad_review"
ignore_cols = [label, "review", "review_clean"]
features = [c for c in reviews_df.columns if c not in ignore_cols]


# split the data into train and test
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(reviews_df[features], reviews_df[label], test_size = 0.20, random
```



Model Training for Classification & Duplicate Column Handling

- 🔍 **Identify Duplicates:** Scrutinize column names to detect any duplications.
- 📄 **Rename Duplicates:** Apply a systematic method to rename duplicate columns, ensuring uniqueness.
- 🚀 **Prepare Data:** Conduct this preprocessing *before* splitting the data for consistency.
- 📊 **Guarantee Compatibility:** Ensure data is in a clean, compatible format for optimal classifier performance.
- ✂️ **Split Dataset:** Divide the preprocessed dataset into training and testing subsets using `train_test_split`.
- 🌲 **Train RandomForest:** Train a `RandomForestClassifier` on the training data, utilizing the features identified.
- 📈 **Feature Importance:** Extract and analyze feature importance scores from the trained RandomForest model to understand feature relevance.

-  **Evaluate Model:** Evaluate model performance using appropriate metrics (e.g., AUC-ROC, Precision-Recall) on the test dataset.

```
In [23]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier

duplicate_columns = reviews_df.columns[reviews_df.columns.duplicated()]

if len(duplicate_columns) > 0:
    print("Duplicate columns found: ", duplicate_columns)
    new_columns = []
    count_dict = {}
    for col in reviews_df.columns:
        if col in count_dict:
            count_dict[col] += 1
            new_columns.append(f"{col}_{count_dict[col]}")
        else:
            count_dict[col] = 0
            new_columns.append(col)
    reviews_df.columns = new_columns

# feature selection
label = "is_bad_review"
ignore_cols = [label, "review", "review_clean"]
features = [c for c in reviews_df.columns if c not in ignore_cols]

# split the data into train and test
X_train, X_test, y_train, y_test = train_test_split(
    reviews_df[features], reviews_df[label], test_size=0.20, random_state=42
)

# Train the Random Forest Classifier
rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train)

# Check lengths and create the DataFrame
print(f"Number of features: {len(X_train.columns)}")
print(f"Number of feature importances: {len(rf.feature_importances_)}")

if len(X_train.columns) == len(rf.feature_importances_):
```

```

feature_importances_df = pd.DataFrame(
    {"feature": X_train.columns, "importance": rf.feature_importances_}
).sort_values("importance", ascending=False)
print(feature_importances_df.head(20))
else:
    print("Error: Feature list and feature importances have different lengths.")
    print("Debugging information:")
    print("X_train columns:", X_train.columns)

```

Number of features: 3832

Number of feature importances: 3832

	feature	importance
3	compound	0.044380
2	pos	0.025604
0	neg	0.024508
3830	total_chars	0.023911
3831	total_words	0.018288
1	neu	0.017590
2846	word_room	0.011517
2232	word_nothing	0.010200
277	word_bad	0.009894
942	word_dirty	0.008723
1937	word_location	0.008266
3195	word_staff	0.007653
3209	word_star	0.007624
1631	word_hotel	0.007175
2277	word_old	0.006407
2196	word_never	0.006181
3081	word_small	0.005772
2510	word_poor	0.005631
424	word_breakfast	0.005502
2860	word_rude	0.005249

In []:

Receiver operating characteristic example

In [24]: *# ROC curve*

```

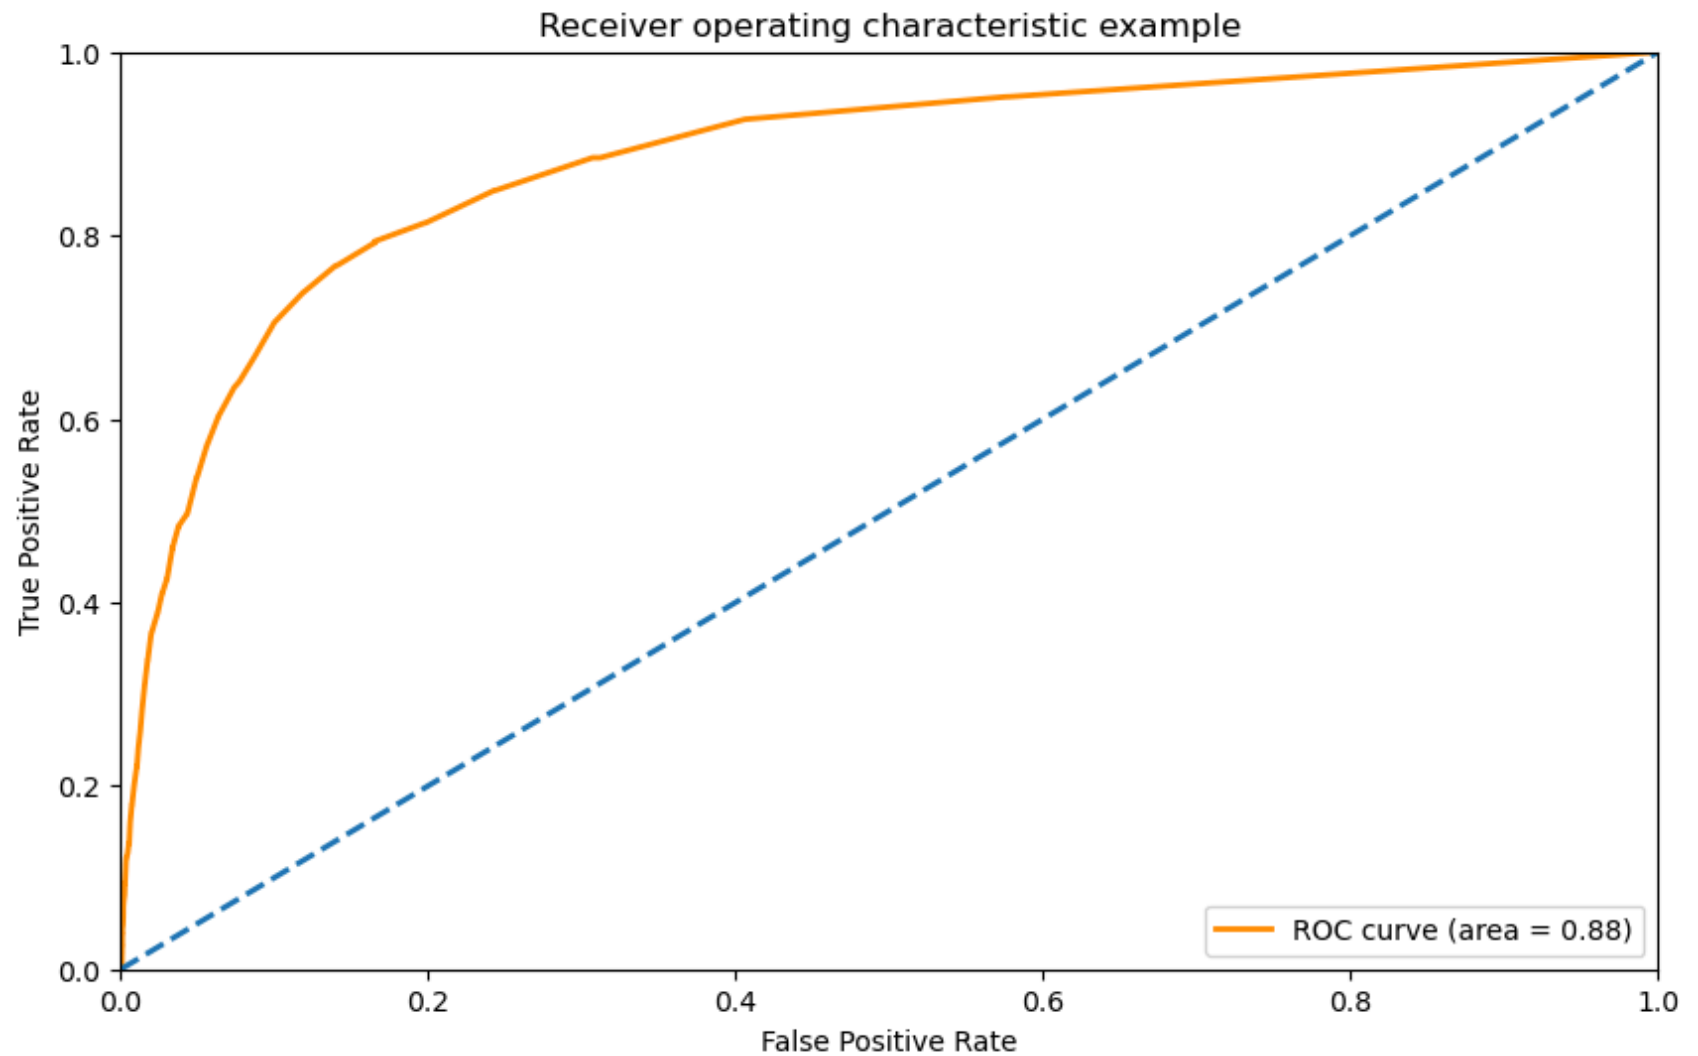
from sklearn.metrics import roc_curve, auc, roc_auc_score
import matplotlib.pyplot as plt

```

```
y_pred = [x[1] for x in rf.predict_proba(X_test)]
fpr, tpr, thresholds = roc_curve(y_test, y_pred, pos_label = 1)

roc_auc = auc(fpr, tpr)

plt.figure(1, figsize = (10, 6))
lw = 2
plt.plot(fpr, tpr, color='darkorange',
         lw=lw, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.0])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic example')
plt.legend(loc="lower right")
plt.show()
```








- **ROC/AUC Evaluation: 0.88**
 - The Receiver Operating Characteristic (ROC) curve and Area Under the Curve (AUC) are standard metrics for binary classifier evaluation.
 - An AUC-ROC of 0.88 indicates strong discriminatory power.
- **Class Imbalance Impact:**
 - Significant class imbalance (high prevalence of negative instances) compromises the ROC curve's reliability in this scenario.

- **False Positive Rate (FPR) Suppression:**
 - The False Positive Rate ($FPR = \text{False Positives} / \text{Negatives}$) is suppressed due to the large number of negative instances.
 - This suppression allows the model to generate numerous false positives while maintaining a low FPR.
- **Artificial AUC Inflation:**
 - The artificially low FPR leads to an inflated AUC-ROC, misrepresenting the model's true performance.
- **Alternative Metrics Recommendation:**
 - Alternative evaluation metrics, robust to class imbalance, are recommended for accurate model assessment.

In []:

Precision-Recall (PR) Curve Analysis

-  **Performance Visualization:** The PR curve visually represents a classifier's performance, especially valuable in imbalanced datasets.
-  **Precision Focus:** Measures the proportion of correctly predicted positive cases out of all predicted positives (minimizes false positives).
- **Recall Focus:** Measures the proportion of correctly predicted positive cases out of all actual positives (minimizes false negatives).
-  **Imbalanced Data Relevance:** Provides a more informative assessment than ROC curves when dealing with significant class imbalances.
-  **Average Precision (AP):** Summarizes the PR curve into a single value, indicating the average precision across different recall thresholds.

In []:

Visualize 2-Class Precision-Recall Curve

```
In [27]: import matplotlib.pyplot as plt
from sklearn.metrics import average_precision_score, precision_recall_curve

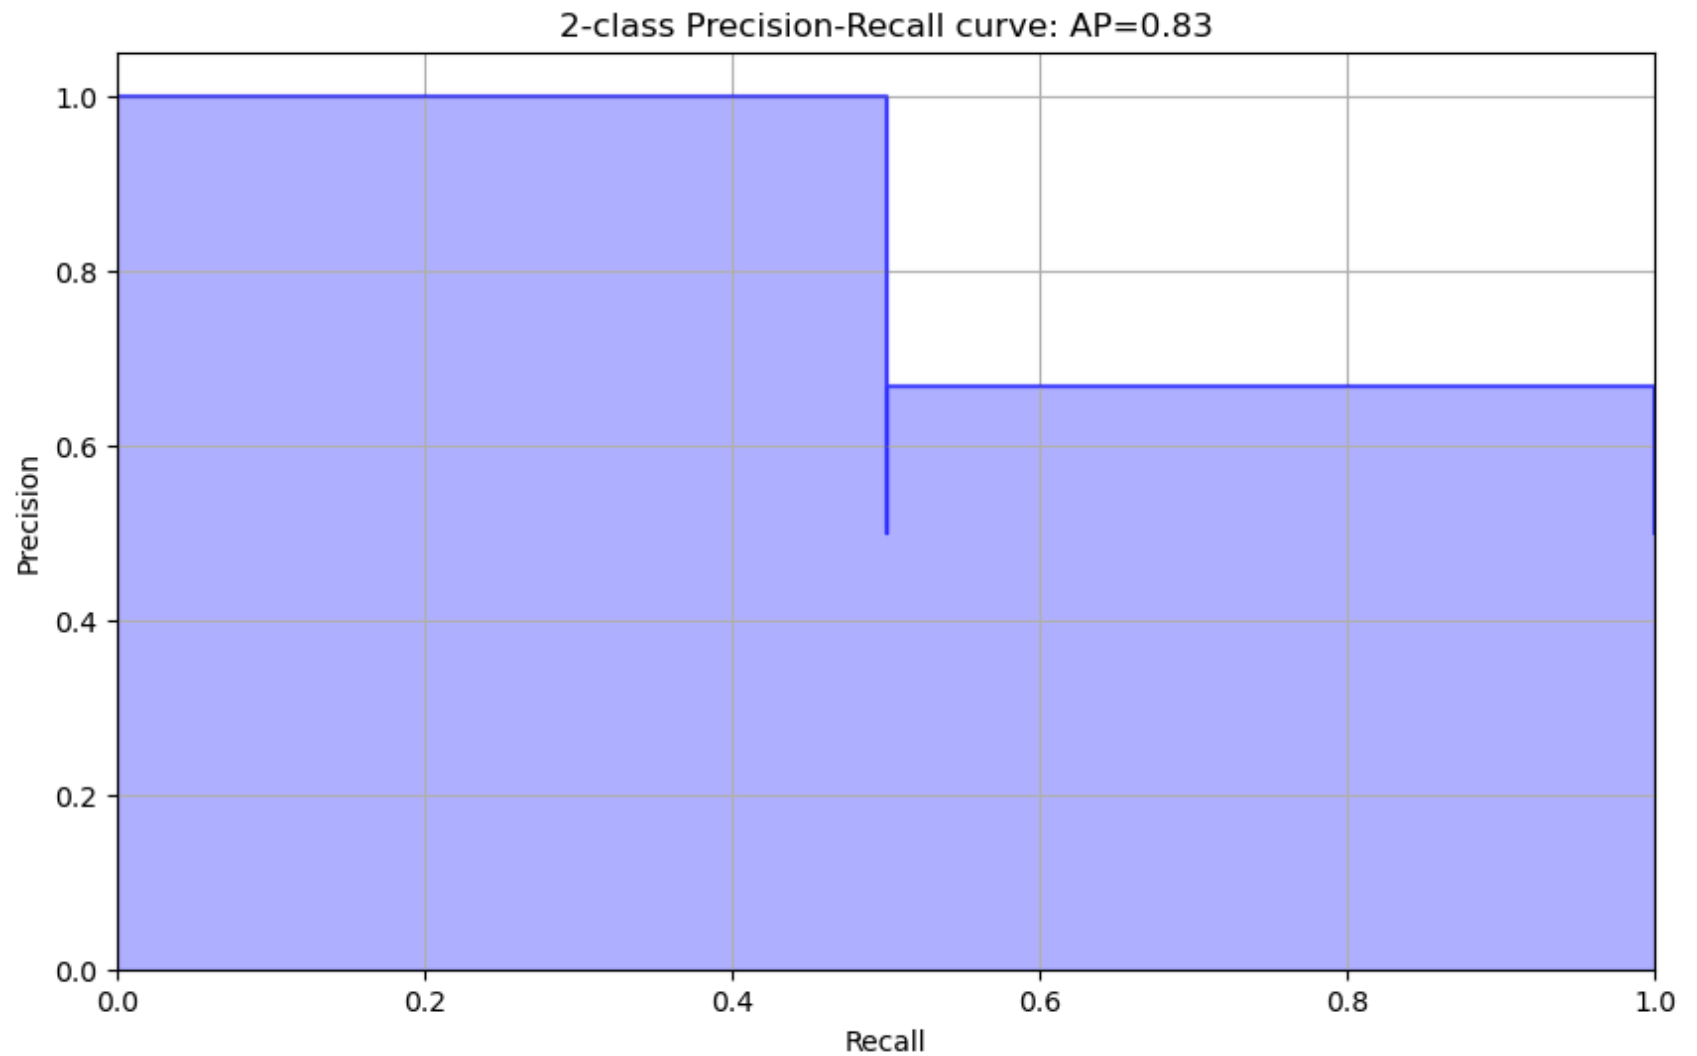
# Compute average precision score
average_precision = average_precision_score(y_test, y_pred)

# Compute Precision-Recall curve
```

```
precision, recall, _ = precision_recall_curve(y_test, y_pred)

# Plot Precision-Recall curve
plt.figure(figsize=(10, 6))
plt.step(recall, precision, color='b', alpha=0.6, where='post')
plt.fill_between(recall, precision, alpha=0.3, color='b', step='post')

plt.xlabel('Recall')
plt.ylabel('Precision')
plt.ylim([0.0, 1.05])
plt.xlim([0.0, 1.0])
plt.title(f'2-class Precision-Recall curve: AP={average_precision:.2f}')
plt.grid(True)
plt.show()
```




Model Performance Metrics

```
In [28]: import matplotlib.pyplot as plt

# Create a Visual Representation of AP and ROC-AUC
fig, ax = plt.subplots(figsize=(6, 3))
ax.bar(['AP', 'ROC-AUC'], [0.83, 0.88], color=['blue', 'green'], alpha=0.7)
ax.set_ylim([0, 1])
```

```
ax.set_ylabel('Score')
ax.set_title('Model Performance Metrics')
plt.grid(axis='y', linestyle='--', alpha=0.6)
plt.show()
```



Model Performance Evaluation

Key Metrics

◆ Precision-Recall Score (PR) = 0.83

- PR summarizes the Precision-Recall curve.
- A higher PR score indicates better performance in handling class imbalance.

◆ ROC-AUC = 0.88

- Measures how well the model distinguishes between classes.
 - A value of 0.88 suggests strong classification ability.
-



Performance Interpretation

- With **PR = 0.83**, the model balances **Precision** and **Recall** effectively.
 - A **ROC-AUC of 0.88** indicates strong **discrimination power** between classes.
 - Generally, a PR score > **0.80** and ROC-AUC > **0.85** represent a well-performing model.
-



Conclusion

- This model demonstrates **high accuracy and reliability**.
 - It is a **strong candidate for deployment** in real-world scenarios.
 - Further tuning may improve results, but current performance is already **robust**.
-



Final Verdict: ★ ★ ★ ★ ★ Good Model!

In []: